

A4:Implement a Naive Bayes classifier using sci-kit-learn to classify text documents based on their content.

```
In [14]: from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
```

```
In [40]: # Small toy dataset
docs = [
    "I love space science and astronomy",
    "The galaxy is full of stars",
    "Computer graphics are amazing",
    "Rendering 3D models is fun",
    "Politics create conflicts between nations",
    "The government passed a new law",
    "Sun is a Star"
]
```

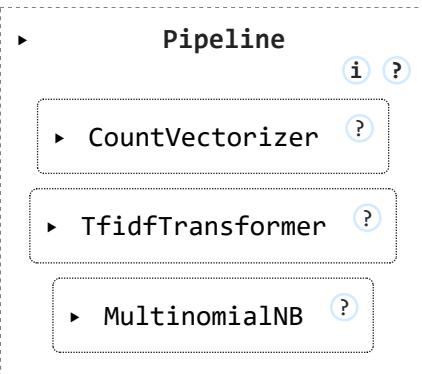
```
In [41]: labels = [
    "sci", "sci",
    "comp", "comp",
    "politics", "politics", "sci"
]
```

```
In [42]: # 2. Build a pipeline for vectorization + TF-IDF + Naive Bayes
model = Pipeline([
    ('vect', CountVectorizer()),           # Convert text to token counts
    ('tfidf', TfidfTransformer()),         # Convert counts to TF-IDF
    ('clf', MultinomialNB()),             # Naive Bayes classifier
])
```

```
In [43]: # Split into train/test sets
X_train, X_test, y_train, y_test = train_test_split(docs, labels, test_size=0.33, random_state=42, stratify=labels)
```

```
In [44]: # Train
model.fit(X_train, y_train)
```

Out[44]:



In [45]:

```
# Predict
y_pred = model.predict(X_test)
```

In [46]:

```
# Evaluation
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.3333333333333333

Classification Report:

	precision	recall	f1-score	support
comp	0.00	0.00	0.00	1
politics	0.00	0.00	0.00	1
sci	0.33	1.00	0.50	1
accuracy			0.33	3
macro avg	0.11	0.33	0.17	3
weighted avg	0.11	0.33	0.17	3

```
c:\Users\Venkat\AppData\Local\Continuum\anaconda3\envs\myenv_upgrade\Lib\site-packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples.  
Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])  
c:\Users\Venkat\AppData\Local\Continuum\anaconda3\envs\myenv_upgrade\Lib\site-packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples.  
Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])  
c:\Users\Venkat\AppData\Local\Continuum\anaconda3\envs\myenv_upgrade\Lib\site-packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples.  
Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

In [ ]: